

Homework 1: Language models (50 points)

Contents

- Logistics
- Exercise 1: Understanding language modeling (12 points)
- Exercise 2: Extracting LLM fingerprints (15 points)
- Exercise 3: Fine-tuning GPT-2 for QA (23 points)

The first homework focuses on the following skills: being able to work with simple formal exercises on language modeling, on understanding and being able to extract properties and configurations of state-of-the-art language models and, finally, training language models yourself!

Logistics

- submission deadline: May 15th 23:59 German time via Moodle
 - please upload a **SINGLE ZIP FILE named Surname_FirstName_HW1.zip** containing the .ipynb file of the notebook (if you solve it on Colab, you can go to File > download), the json file for Ex. 2 and a .png or .jpg file with your losses plot from Ex. 3.
- please solve and submit the homework **individually!**
- if you use Colab, to speed up the execution of the code on Colab (especially Exercise 3), you can use the available GPU (if Colab resources allow). For that, before executing your code, navigate to Runtime > Change runtime type > GPU > Save.

Exercise 1: Understanding language modeling (12 points)

Please answer the following exercises. Importantly, please reason step by step; i.e., where calculations are required, please provide intermediate steps of how you arrived at your solution. You do not need to write any code, just mathematical solutions.

1. [6pts] Consider the corpus C with the following sentences: $C = \{\text{"The cat sleeps", "The mouse sings", "The cat sleeps", "A dog sings"}\}$. (a) Define the vocabulary V of this corpus (assuming by-word tokenization). (b) Pick one of the four sentences in C . Formulate the probability of that sentence in the form of the chain rule. Calculate the probability of each term in the chain rule, given the corpus.
2. [4pts] We want to train a neural network that takes as input two numbers x_1, x_2 , passes them through three hidden linear layers, each with 13 neurons, each followed by the ReLU activation function, and outputs three numbers y_1, y_2, y_3 . Write down all weight matrices of this network with their dimensions. (Example: if one weight matrix has the dimensions 3×5 , write $M_1 \in \mathbb{R}^{3 \times 5}$)
3. [2pts] Consider the sequence: "Input: Some students trained each language model". Assuming that each word+space/punctuation corresponds to one token, consider the following token probabilities of this sequence under some trained language model: $p = [0.67, 0.91, 0.83, 0.40, 0.29, 0.58, 0.75]$. Compute the average surprisal of this sequence under that language model. [Note: in this class we always assume the base e for \log , unless indicated otherwise. This is also usually the case throughout NLP.]

Exercise 2: Extracting LLM fingerprints (15 points)

For this task, your job is to extract the "fingerprint" of a state-of-the-art large language model from the paper. Specifically, your job is to:

- find the model that is assigned to your surname in the list **HW1_Model2Group_assignment.csv** (to be found on Moodle under topic 02). Please investigate the latest version of your model, unless the version is specified in the list.
- find out the following characteristics of your model
- submit a json file with your responses in the following format (below is a partial example).

Note that, of course, it might be that some information is not available or that some categories are not applicable. The idea is, that, as a course we can create a fun website which will show a somewhat comprehensive graphical comparison of current language models and their configurations. Based on your collective json files, the lecturers will set up a front end at some point during the class.

IMPORTANT: Please email the lecturers by the homework deadline if you DO NOT consent that your json file is used for this idea.

```
{
  "model_name": "GPT-35",
  "huggingface_model_id": "gpt35",
  "paper_url": "https://arxiv.org/abs/XXX",
  "tokenizer_type": "BPE",
  "vocabulary_size": "XXX",
  "architecture": "Mixture of transformer agents",
  "architecture_type": "decoder only",
  "architecture_quirks": [
    "sparse attention",
    "...",
  ],
  "parameters": "XXX",
  "finetuning_type": "RLHF",
  "training_data_cutoff": "2050",
  "number_training_tokens": "XXX",
  "pretraining_data_size": "1GB",
  "finetuning_data_size": "XXX",
  "training_data": [
    "Books corpus",
    "Twitter",
    "...",
  ],
  "finetuning_data": [
    "XXX",
    "XXX",
    "...",
  ],
  "access": "open",
  "summary": "A few sentences of what the model claims to be their unique selling point / main contrib
}
```

Exercise 3: Fine-tuning GPT-2 for QA (23 points)

The learning goal of this exercise is to practice fine-tuning a pretrained LM, GPT-2 small, for a particular task, namely commonsense question answering (QA). We will use a task-specific dataset, [CommonsenseQA](#), that was introduced by [Talmor et al. \(2018\)](#). We will evaluate the performance of the model on our test split of the dataset over training to monitor whether the model's performance is improving and compare the performance of the base pretrained GPT-2 and the fine-tuned model. We will need to perform the following steps:

1. Prepare data according to steps described in [sheet 1.1](#)
 1. additionally to these steps, prepare a custom Dataset (like in [sheet 2.3](#)) that massages the dataset from the format that it is shipped in on HuggingFace into strings that can be used for training. Some of the processing steps will happen in the Dataset.
2. Load the pretrained GPT-2 model
3. Set up training pipeline according to steps described in sheet 2.5
4. Run the training while tracking the losses
5. Save plot of losses for submission

Your tasks:

1. [19pts] Complete the code in the spots where there is a comment "#### YOUR CODE HERE ####". There are instructions in the comments as to what the code should implement. With you completed code, you should be

necessarily expect great performance of the fine-tuned model (and the actual performance will *not* be graded).

Often there are several correct ways of implementing something. Anything that is correct will be accepted.

2. [4pts] Answer questions at the end of the exercise.

```
# note: if you are on Colab, you might need to install some requirements
# as we did in Sheet 1.1. Otherwise, don't forget to activate your local environment

from datasets import load_dataset
from transformers import AutoTokenizer, AutoModelForCausalLM, GPT2Tokenizer, GPT2LMHeadModel
import torch
from torch.utils.data import DataLoader
from torch.utils.data import Dataset
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm
```

```
# additionally, we need to install accelerate
# uncomment and run the following line on Colab or in your environment
# !pip install accelerate
# NOTE: in a notebook, reloading of the kernel might be required after installation if you get dependenc
```

```
### 1. Prepare data with data prepping steps from sheet 1.1

# a. Acquiring data
# b. (minimally) exploring dataset
# c. cleaning / wrangling data (combines step 4 from sheet 1.1 and step 1.1 above)
# d. splitting data into training and test set (we will not do any hyperparam tuning)
# (we don't need further training set wrangling)
# e. tokenizing data and making sure it can be batched (i.e., converted into 2d tensors)
# this will also happen in our custom Dataset class (common practice when working with text data)
```

```
# download dataset from HF
dataset = load_dataset("tau/commonsense_qa")
```

```
# inspect dataset
print(dataset.keys())
# print a sample from the dataset
### YOUR CODE HERE ###
```

Note that the test split does not have ground truth answer labels. Therefore, we will use the validation split as our test split.

```
# load tokenizer
tokenizer = AutoTokenizer.from_pretrained("gpt2")
tokenizer.pad_token = tokenizer.eos_token
# set padding side to be left because we are doing causal LM
tokenizer.padding_side = "left"
```

```
def message_input_text(example):
    """
    Helper for converting input examples which have
    a separate question, labels, answer options
    into a single string.

    Arguments
    -----
    example: dict
        Sample input from the dataset which contains the
        question, answer labels (e.g. A, B, C, D),
        the answer options for the question, and which
        of the answers is correct.

    Returns
    -----
    input_text: str
        Formatted training text which contains the question,
        the formatted answer options (e.g., 'A. <option 1> B. <option 2>' etc)
        and the ground truth answer.
    """
    # combine each label with its corresponding text
    answer_options_list = list(zip(
        example["choices"]["label"],
```

```

# join each label and text with . and space
answer_options = ### YOUR CODE HERE ###
# join the list of options with spaces into single string
answer_options_string = ### YOUR CODE HERE ###
# combine question and answer options
input_text = example["question"] + " " + answer_options_string
# append the true answer with a new line, "Answer: " and the label
input_text += "\nAnswer: " + example["answerKey"]

return input_text

# process input texts of train and test sets
massaged_datasets = dataset.map(
    lambda example: {
        "text": message_input_text(example)
    }
)

```

```

# inspect a sample from our preprocessed data
massaged_datasets["train"][0]

```

```

class CommonsenseQADataset(Dataset):
    """
    Custom dataset class for CommonsenseQA dataset.
    """

    def __init__(
        self,
        train_split,
        test_split,
        tokenizer,
        max_length=64,
    ) -> None:
        """
        Initialize the dataset object.

        Arguments
        -----
        train_split: dict
            Training data dictionary with different columns.
        test_split: dict
            Test data dictionary with different columns.
        tokenizer: Tokenizer
            Initialized tokenizer for processing samples.
        max_length: int
            Maximal length of inputs. All inputs will be
            truncated or padded to this length.
        """
        self.train_split = train_split['text']
        self.test_split = test_split['text']
        self.tokenizer = tokenizer
        self.max_length = max_length

    def __len__(self):
        """
        Method returning the length of the training dataset.
        """
        return ### YOUR CODE HERE ###

    def __getitem__(self, idx):
        """
        Method returning a single training example.
        Note that it also tokenizes, truncates or pads the input text.
        Further, it creates a mask tensor for the input text which
        is used for causal masking in the transformer model.

        Arguments
        -----
        idx: int
            Index of training sample to be retrieved from the data.

        Returns
        -----
        tokenized_input: dict
            Dictionary with input_ids (torch.Tensor) and an attention_mask
            (torch.Tensor).
        """
        # retrieve a training sample at the specified index idx
        input_text = ### YOUR CODE HERE ###
        tokenized_input = self.tokenizer(
            input_text,
            max_length=### YOUR CODE HERE ###
            padding="max_length"
        )

```

```

        return_tensors="pt"
    )
    tokenized_input["attention_mask"] = (tokenized_input["input_ids"] != tokenizer.pad_token_id).long()
    return tokenized_input

```

```

# move to accelerated device
if torch.cuda.is_available():
    device = torch.device("cuda")
    print(f"Device: {device}")
elif torch.backends.mps.is_available():
    device = torch.device("mps")
    print(f"Device: {device}")
else:
    device = torch.device("cpu")
    print(f"Device: {device}")

```

```

# 2. init model

# load pretrained gpt2 for HF
model = ### YOUR CODE HERE ###
# print num of trainable parameters
model_size = sum(t.numel() for t in model.parameters())
print(f"GPT-2 size: {model_size/1000*2:.1f}M parameters")

```

Hint: If you run out of memory while trying to run the training, try decreasing the batch size.

```

# 3. set up configurations required for the training loop

# instantiate dataset with the downloaded commonsense_qa data
train_dataset = CommonsenseQADataset(
    ### YOUR CODE HERE ###
)
# create a DataLoader for the dataset
# the data loader will automatically batch the data
# and iteratively return training examples (question answer pairs) in batches
data_loader = DataLoader(
    train_dataset,
    batch_size=32,
    shuffle=True
)

```

```

# 4. run the training of the model
# Hint: for implementing the forward pass and loss computation, carefully look at the exercise sheets
# and the links to examples in HF tutorials.

# put the model in training mode
model.train()
# move the model to the device (e.g. GPU)
model = model.to(device)

# training configurations
# feel free to play around with these
epochs = 1
train_steps = len(train_dataset) // 32
print("Number of training steps: ", train_steps)
# define optimizer and learning rate
optimizer = torch.optim.AdamW(model.parameters(), lr=5e-4)
# define some variables to accumulate the losses
losses = []

# iterate over epochs
for e in range(epochs):
    # iterate over training steps
    for i in tqdm(range(train_steps)):
        # get a batch of data
        x = next(iter(data_loader))
        # move the data to the device (GPU)
        x = ### YOUR CODE HERE ###

        # forward pass through the model
        ### YOUR CODE HERE ###
        outputs = model(
            ### YOUR CODE HERE ###
        )
        # get the loss
        loss = ### YOUR CODE HERE ###
        # backward pass
        ### YOUR CODE HERE ###

```

```

    ### YOUR CODE HERE ###

    # zero out gradient for next step
    ### YOUR CODE HERE ####

    # evaluate on test set every 10 steps
    if i % 10 == 0:
        print(f"Epoch {e}, step {i}, loss {loss.item()}")
        # TODO

```

5. Plot the fine-tuning loss and MAKE SURE TO SAVE IT AND SUBMIT IT

```

# plot training losses on x axis
plt.plot(### YOUR CODE HERE ####)
plt.xlabel("Training steps")
plt.ylabel("Loss")

```

print a few predictions on the eval dataset to see what the model predicts

construct a list of questions without the ground truth label
and compare prediction of the model with the ground truth

```

def construct_test_samples(example):
    """
    Helper for converting input examples which have
    a separate question, labels, answer options
    into a single string for testing the model.

    Arguments
    -----
    example: dict
        Sample input from the dataset which contains the
        question, answer labels (e.g. A, B, C, D),
        the answer options for the question, and which
        of the answers is correct.

    Returns
    -----
    input_text: str, str
        Tuple: Formatted test text which contains the question,
        the formatted answer options (e.g., 'A. <option 1> B. <option 2>' etc);
        the ground truth answer label only.
    """

    answer_options_list = list(zip(
        example["choices"]["label"],
        example["choices"]["text"]
    ))
    # join each label and text with . and space
    answer_options = ### YOUR CODE HERE ####
    # join the list of options with spaces into single string
    answer_options_string = ### YOUR CODE HERE ####
    # combine question and answer options
    input_text = example["question"] + " " + answer_options_string
    # create the test input text which should be:
    # the input text, followed by the string "Answer: "
    # we don't need to append the ground truth answer since we are creating test inputs
    # and the answer should be predicted.
    input_text += ### YOUR CODE HERE ####

    return input_text, example["answerKey"]

test_samples = [construct_test_samples(dataset["validation"][i]) for i in range(10)]
test_samples

```

Test the model

```

# set it to evaluation mode
model.eval()
model = model.to("cpu")

predictions = []
for sample in test_samples:
    input_text = sample[0]
    input_ids = tokenizer(input_text, return_tensors="pt").to(device)
    output = model.generate(
        input_ids=input_ids,
        attention_mask = input_ids.attention_mask,
        max_new_tokens=2,
        do_sample=True,
        temperature=0.4
    )

```

```
predictions.append((input_text, prediction, sample[1]))  
  
print("Predictions of trained model ", predictions)
```

Questions:

1. Provide a brief description of the CommonsenseQA dataset. What kind of task was it developed for, what do the single columns contain?
2. What loss function is computed for this training? Provide the name of the function (conceptual, not necessarily the name of a function in the code).
3. Given your loss curve, do you think your model will perform well on answering common sense questions? (Note: there is no single right answer; you need to interpret your specific plot)
4. Inspect the predictions above. On how many test questions did the model predict the right answer? Compute the accuracy.

Previous

< [Sheet 2.5: Introduction to HuggingFace & LMs](#)